

A4: Computer Networking (II)

Computer Systems 2022
Department of Computer Science
University of Copenhagen

David Marchant

Due: Sunday, 4th of December, 16:00
Version 1 (November 23, 2022)

This is the fifth assignment in the course on Computer Systems 2022 at DIKU and the second on the topic of Computer Networks. We encourage pair programming, so please form groups of 2 or 3 students. Groups cannot be larger than 3, and we strongly recommend that you do not work alone.

For this assignment you will receive a mark out of 4 points; You must attain at least half of the possible points to be admitted to the exam. For details, see the *Course description* in the course page. This assignment belongs to the CN category together with A3. Resubmission is not possible.

It is important to note that only solving the theoretical part will result in 0 points. Thus, the theoretical part can improve your score, but you *have* to give a shot at the programming task.

The web is more a social creation than a technical one. I designed it for a social effect — to help people work together — and not as a technical toy. The ultimate goal of the Web is to support and improve our weblike existence in the world. We clump into families, associations, and companies. We develop trust across the miles and distrust around the corner.

— Tim Berners-Lee, *Weaving the Web* (1999)

Overview

This assignment has two parts, namely a theoretical part (Section 1) and a programming part (Section 2). The theoretical part deals with questions that have been covered by the lectures. The programming part requires you to complete a file transfer service using socket programming in C. This assignment uses a similar, but not identical, protocol to A3 in order to transfer files according to a peer-to-peer architecture. In this assignment, the programming effort relies on building a peer consisting of concurrent client and server interactions. More details will follow in the programming part (Section 2).

1 Theoretical Part (25%)

Each section contains a number of questions that should be answered **briefly** and **precisely**. Most of the questions can be answered within 2 sentences or less. Annotations have been added to questions that demand longer answers, or figures with a proposed answer format. Miscalculations are more likely to be accepted, if you account for your calculations in your answers.

1.1 Transport protocols

The questions relating to TCP in this section do not assume knowledge of the flow and congestion control parts of TCP.

1.1.1 TCP reliability and utilization

Part 1: Is the 3-way-handshake in TCP really needed? Can you suffice with less? Explain why or why not.

Part 2: How does TCP facilitate a *full-duplex*¹ connection?

1.1.2 Reliability vs overhead

Part 1: Considering the data transferred over a network, how does a TCP connection add overhead relative to UDP?

Part 2: Explain the TCP notion of reliable data transfer. In which way are TCP connections reliable? Are packets **always** delivered?

1.2 TCP: Principles and practice

Some of the questions below refer to one or more *Request For Comments* (RFC) memorandums posted by the *Internet Engineering Task Force* (IETF). In most cases, you can find answers to the question in the book, but we recommend that you browse through these RFCs to see how internet standards are formulated and distributed in practice.

Be careful when answering the questions below. If any ambiguity may arise, remember to specify the point of view of your answers – is it from the server or the client or both?

1.2.1 TCP headers

Part 1: The TCP segment structure is shown in figure 3.29 in K&R (or section 1.3 of RFC793²).

- 1.1. What is the purpose of the RST bit? Give an example of when a TCP stack may return a RST packet.

¹Full speed bidirectional data transfer

²<https://tools.ietf.org/html/rfc793#section-3.1>

- 1.2. What does the sequence number and acknowledgement number headers refer to for a given connection? Is there any relation between server and client acknowledgement and sequence numbers?
- 1.3. What is the purpose of the window? What does a positive window size signify?
- 1.4. If the window size of a TCP receiver is 0, what happens then at the sender? How can the sender find out when the receiver window is not 0?

Part 2: In one of the previous questions, you were asked why the 3-way handshake was necessary. One of the practical reasons for this 3-way-handshake was the synchronization of sequence numbers. In some situations, the client can initiate its transfer of data after having received the SYN-ACK-packet. Explain when this is possible. (*Hint: What does the server need to tell the client in the SYN-ACK-packet?*)

1.2.2 Flow and Congestion Control

Part 1: Explain the type of congestion control that modern TCP implementations employ. Is it network assisted congestion control?

Part 2: Congestion control translates directly to limiting the transmission rate. How does TCP determine the transmission rate when doing congestion control?

Part 3: How does a TCP sender infer the value of the congestion window, that is, what is the basis for calculating the congestion window? (*Answer with at most 10 sentences*)

Part 4: The fast retransmit extension, as described in RFC 2581³ allows for retransmission of packets after receiving 3 duplicate ACKs for the same segment. What does the sender need to implement in order to allow for fast retransmit? What does the receiver need to implement? If a sender supports fast retransmit and needs to do fast retransmit for a specific segment, how many ACKs should the sender receive for the previous packet?

³<https://tools.ietf.org/html/rfc2581>

2 Programming Part (75 %)

For the programming part of this assignment, you will implement a peer in a peer-to-peer file-sharing service. Many of the programming tasks set in the exercises relate to the sub-tasks within this assignment, so if you are stuck at any point consider revisiting them.

2.1 Design and overview

The file-sharing service consists of many identical peer processes. When a peer starts it must be given the address of a peer already on the network, with whom it will make contact to join the network. The first peer to join will not attempt to connect to anyone, as there is no one yet to connect to. This design uses a custom protocol outlined below.

In this protocol each peer can perform either client or server interactions, and must be capable of performing each concurrently. Note that this may also mean that different server interactions may be performed concurrently.

The first peer to join a network cannot perform client interactions, as there would be no peers to interact with. Therefore it will only perform server interactions. We shall refer to this as the *initial peer*, though do note that it is programmatically identical to any other peer. If a new peer, referred to here as the *joining peer* were to attempt to join the network it would first contact the *initial peer* to make itself known to the network. The *initial peer* will then update its record of all peers on the network and reply to the *joining peer* with a complete list of all peers on the network. The *initial peer* will then inform any other peers on the network that the *joining peer* has joined, so that all peers will have an up to date record of all peers on the network. If a third peer or fourth peer joined they could do so in the same way, using any peer already on the network as their initial peer to contact.

Note that in this assignment we are not overly concerned with security, or if people are allowed to access files. What we are concerned by is the ability to construct a network node (the peer) in such a way that it can act as a client and server at the same time without deadlocking, or causing race conditions.

Note that as in A3, all messages are limited in how many bytes can be sent at once. If a peer requests a file larger than this limit then it will therefore be sent over several messages which the peer will have to reassemble.

For this assignment you are going to implement a peer that registers with a network, retrieves files, and responds appropriately to those same requests. You can find a complete description of the protocol in Section 3.

2.2 API and Functionality

2.2.1 Key Implementation Tasks

The peer you will be implementing needs to perform the six basic tasks listed below. Each should be completed according to the protocol described in Section 3. The client should:

- 2.1. Registering a new peer on the network. The joining peer will send its IP and port to a known peer already on the network, and receive a complete list of the network as a reply.
- 2.2. Respond to peer requests to join the network. The peer should add the requesting contact details to its record of the network, and respond with a complete account of the network.
- 2.3. Retrieve a small file. The user should request a file small enough to fit into a single message. For this the example file `tiny.txt` has been provided, though you are free to test on others. This should be retrieved by the peer and written into its own file storage. The client should validate the received file to check that it has been received correctly.
- 2.4. Respond to requests for files small enough to fit into a single message. For this the example file `tiny.txt` has been provided, though you are free to test on others.
- 2.5. Retrieve a large file. The user should request a file large enough to require several messages to completely send. For this the example file `hamlet.txt` has been provided, though you are free to test on others. This should be retrieved by the peer and written into its own file storage. The client should validate the received file to check that it has been received correctly. Note that you should not assume that blocks are received in order.
- 2.6. Respond to requests for files large enough to require several messages to completely send. For this the example file `hamlet.txt` has been provided, though you are free to test on others. Note that the protocol makes no requirement of the order in which blocks are sent, e.g. it is up to your implementation to make a sensible choice.

2.2.2 Test environment

One of the difficulties with developing a system that relies on network communication is that besides a network, you need to have two or more peers running, and potentially debug both simultaneously.

To assist in developing and testing your code, we have provided an implementation of the peer, but written in Python. You can use this to get started with a setup that runs fully contained on your own machine. By running a local Python peer and your own C peer it is possible to debug and monitor all pieces of the communication, which is often not possible with an existing system.

Note that there are two directories within the Python directory, *first_peer* and *second_peer*. This is as you should be testing two sides of interaction with your C peer. E.g. test that your C peer can register with the *first_peer* and retrieve files from it. Also test from the *second_peer* that it can request from your C Peer and have files sent to it when requested. You should also test with all 3 at the same time so as to test your concurrency setup as you CANNOT assume that all client interactions will complete before all server interactions start or vice versa. A sleep function on line 408 of the Python peer has been suggested

to help aide testing these interactions. Note that both Python peers are identical (technically they should be literally the same file just linked), though have been given different config.yaml files and available data files by default. Feel free to alter these configs and files as you wish, they are merely suggestions for some configurations to test.

You can of course also peek into or alter the Python and get inspiration for implementing you own peer, but beware that what is a sensible design choice in Python *may* not be a good choice for C. You are also reminded that although you can alter the Python as much as you want to provide additional debugging or the like, they are currently correct implementations of the defined protocols. Be careful in making changes that you do not alter the implemenation details, as this may lead your C implementation astray.

To start an appropriate test environment you can run a peer using the command below, run from the *python/first_peer* directory:

```
python3 peer.py config.yaml
```

Assuming you have not altered *python/first_peer/config.yaml*, this will start a peer at 127.0.0.1:12345. This will be capable of serving any files within the *python/first_peer* directory. You can now start another peer to connect to this one. You could use either a second Python peer within *python/second_peer*, or the C. You may need to alter the default configs to get the functionality you require. Note that previously we have used 0.0.0.0 as our address, which tells your system not to use your network driver at all for this connection. Some students have reported this hasn't worked on their machines so we have switched to 127.0.0.1, which will tell your system to use the network driver but route the connection back to the host. Though in practice this means that both are doing essentially the same thing, if you find one does not work try the other, and if that doesn't work talk to a TA.

To make and deploy the C peer run the following within the *src* directory:

```
make  
./peer config.yaml
```

By default this will compile your client, though as not yet implemented it will not do anything significant.

Do note that both the Python and C peers as handed out act as simple linear scripts for their client interactions, registering a peer and then retrieving the two required files. This structure has been picked for ease of development and understandability only, and is not intended as demonstration of good practice. If you would rather alter this structure so as to allow for a peer to dynamically request whatever files they wish whenever they wish, then do feel free to do so, though it is not aim or requirement of this assignment.

2.3 Run-down of handed-out code and what is missing

- *src/peer.c* contains the peer, and is where it is expected all of your coding will take place. If you alter any other files, make sure to highlight this in your report. This code contains five guide functions, *server_thread*,

`handle_server_request`, `handle_register`, `handle_inform` as well as `handle_retrieve`. These should illustrate where you might begin, but feel free to make whatever alterations to the structure you wish. Your edits WILL NOT be limited to these functions however, and you will NEED to make additional edits to the `send_message` function and potentially others. The `send_message` demonstrates functionality from A3 so you *may* wish to replace it with your own solution to that if you are more familiar with it, though note it does not encompass all of the functionality expected from A4. Your completed system should write any retrieved files to the `src` directory.

- `src/peer.h` contains a number of structs you may find useful in building your client. You are not required to use them as is, or at all if you would rather solve the problem some other way.
- `src/config.yaml` contains the definitions for a peer to be hosted, acting as a joining peer in the network.
- `src/sha256.c` and `src/sha256.h` contain an open-source stand-alone implementation of the SHA256 algorithm, as used throughout the network. You should not need to alter this file in any way.
- `src/common.c` and `src/common.h` contains some helper functions to assist in parsing the configuration file. You should not need to alter this file in any way.
- `src/csapp.c` and `src/csapp.h` contains definitions for robust read and write operations. You should not need to alter this file in any way.
- `src/endian.h` contains definitions for converting between big and small endian numbers. This is only useful on a mac, and should not be necessary in the majority of solutions so do not be concerned if you do not use any of these functions. You should not need to alter this file in any way.
- `python/first_peer/peer.py` and `python/second_peer/peer.py` contains the peer, written in Python. This can be run on Python 3.7 or newer, and requires a config file written in the YAML format. You should not need to alter this file, though may find adding more debug print statements helpful.
- `python/first_peer/config.yaml` contains the definitions for a peer to be hosted, acting as an initial peer in the network.
- `python/second_peer/config.yaml` contains the definitions for a peer to be hosted, acting as an joining peer in the network.
- `python/first_peer/tiny.txt` and `python/first_peer/hamlet.txt` are example data files which if you can transfer to and from the peer correctly are enough to say you have met some of the goals of this assignment. Testing using additional files would be advantageous in a report.
- README contains the various commands you will need to get this project up and running.

2.4 Testing

For this assignment, it is *not* required of you to write formal, automated tests, but you *should* test your implementation to such a degree that you can justifiably convince yourself (and thus the reader of your report) that each API functionality implemented works, and are able to document those which do not.

Simply running the program, emulating regular user behaviour and making sure to verify the result file should suffice, but remember to note your results. When testing parallel systems that may act non-deterministically any amount of user testing is of limited utility. However, you should attempt to stress-test your system say by running many peers at once, or by maximising chances of race conditions.

One final resource that has been provided to assist you is that an instance of `peer.py` has been remotely hosted, and which you can connect to. This is functionally identical to those contained in the handout, with the obvious difference that any communications with it will be properly over a network.

The server is designed to be robust, so should be resilient to any malformed messages you send, but as it is only a single small resource be mindful of swamping it with requests and only use them once you are confident in your system. You can connect to the peer at the following address and port:

Test Server: 130.225.104.138:5555

Note that depending on whatever firewalls or other network configuration options you have, you may not be able to reach the server from home, but that you should be able to from within the university. Additionally, note that this version of the Peer will slightly differ from the presented Python peer, in that it will automatically remove peers from its network it has not heard from for 5 mins. This will prevent old test trackers from clogging up the network and is not required functionality for you to emulate.

2.5 Recommended implementation progress

As mentioned in the previous section, `src/peer.c` contains five unimplemented functions and at least one partially completed function. The satisfaction of these should yield a functioning peer. Do note that you are not limited to completing these functions and are expected to make changes even through the provided functions so as to ensure functionality that avoids deadlock AND race conditions. It is expected you will stick to modifying `peer.c`, and should specifically highlight in your report if you deviate from this.

In this section, we give a short recap of your implementation todos; this can serve as a checklist for your project, and *we recommend* that you work on them in the order presented here.

- 2.1. With the peer acting as a client, request to register with another peer and appropriately manage any feedback
- 2.2. With the peer acting as a client, request to retrieve both the small and large file from another peer and appropriately manage any feedback

- 2.3. With the peer acting as a server, respond to inform requests from another peer
- 2.4. With the peer acting as a server, respond to register requests from another peer and generate appropriate feedback
- 2.5. With the peer acting as a server, respond to retrieve requests from another peer and generate appropriate feedback
- 2.6. Throughout these tasks ensure that your peer can act as a client and server AT THE SAME TIME. You must ensure that any race conditions or deadlocks are appropriately handled
- 2.7. Manually test your implementation, documenting bugs found and how you fix them (if you are able to).
- 2.8. Meanwhile, do not neglect the theoretical questions. They may be relevant to your understanding of the implementation task.

2.6 Report and approximate weight

The following approximate weight sums to 75% and includes the implementation when relevant.

Please include the following points in your report:

- Discuss the provided protocol and how you were able to ensure a thread-safe implementation. For instance, what data needed to be shared between the client and server side of the application, and how did you avoid race conditions? Have you done so in a way that also avoids deadlocks? What circumstances would stress your implementation either by creating errors, or by degrading performance? (Approx. weight: 15%)
- Document technical implementation you made for the peer - cover in short each of the six tasks (as client register, as client retrieve, as server register, as server inform, as server retrieve), as well as any additional changes you made. Each change made should be briefly justified. You should also describe what input/data your implementation is capable of processing and what it is not. (Approx. weight: 30%)
- Discuss how your design was tested. What data did you use on what machines? Did you automate your testing in any way? Remember that you are not expected to have built a perfect solution that can manage any and all input, but you are always meant to be able to recognise what will break your solution. Testing on the provided two files (tiny.txt and hamlet.txt) is not sufficient, you should also test using data that does not work, or with requests that are incorrect. You should also test with more than 2 peers running at once. (Approx. weight: 15%)
- Discuss any shortcomings of your implementation, and how these might be fixed. It is not necessarily expected of you to build a fully functional

peer, but it *is* expected of you to reflect on the project. Even if you implement the protocol exactly as described throughout this document there will still be problems to identify in terms of usability, concurrency or functionality. (Approx. weight: 15%)

As always, remember that it is also important to document half-finished work. Remember to provide your solutions to the theoretical questions in the report pdf.

3 Protocol description

This section defines the implementation details of the components used in the A4 peer-to-peer file-sharing network.

3.1 File structures of the peer

All peers will serve files relative to their own location.

Both the server and client use a base data directory, defined in `config.yaml`. Any files within the servers base directory can be retrieved. Any files retrieved by the client will be placed within its file base.

3.2 Format of the peer client requests

All client messages from the peer must contain the same header as follows:

- 16 bytes - The IP address the sending peer is listening on,
as UTF-8 encoded bytes
- 4 bytes - The port the sending peer is listening on,
unsigned integer in network byte-order
- 4 bytes - The command code describing the nature of this
message, unsigned integer in network byte-order
- 4 bytes - The length of the request body, unsigned integer
in network byte-order

The command codes specified above can only be one of 3 values. These, along with the appropriate request bodies for them are:

- 1 - Register, when a peer is attempting to make first
contact with a network. No body is provided with this
command
- 2 - Retrieve, when a peer is requesting a data file to be
transferred to it from another peer. The request body
must be the filename/filepath of the requested file
- 3 - Inform, when a peer is informing another peer of a
third peer that has just joined the network, The
request body must fit the 20 bytes described below:
 - 16 bytes - The IP address of the newly joined peer,
as UTF-8 encoded bytes
 - 4 bytes - The port of the newly joined peer, unsigned
integer in network byte-order

In the case of the Register and Retrieve (1 2) commands, or if the command could not be determined, a reply will ALWAYS be expected from the other peer. ONLY in the case of the inform (3) command do we not expect a reply.

3.3 Format of the client server responses

All replies from the server to the client must contain the same header as follows:

- 4 bytes - Length of response body,
 unsigned integer in network byte-order
- 4 bytes - Status Code of the response,
 unsigned integer in network byte-order
- 4 bytes - Block number, a zero-based count of which block in
 a potential series of replies this is,
 unsigned integer in network byte-order
- 4 bytes - Block count, the total number of blocks to be sent,
 unsigned integer in network byte-order
- 32 bytes - Block hash, a hash of the response data in this
 message only, as UTF-8 encoded bytes
- 32 bytes - Total hash, a hash of the total data to be sent
 across all blocks, as UTF-8 encoded bytes

In addition to the header each response will include an additional payload of the length given in the header. Note that in the case of a small enough response to only require a single reply, then the block hash and total hash will be identical. If the response is to a request for a data file, and it could be retrieved with no errors then the payload will be either all or part of said file, depending on the size of the file. If the response is to registering on the network, the response will be a list of all known peers on the network. Each entry in the list will be formatted as:

- 16 bytes - The IP address of a network peer,
 as UTF-8 encoded bytes
- 4 bytes - The port of a network peer,
 unsigned integer in network byte-order

The list itself will be some multiple of 20 bytes long, with the first 20 bytes being one peer, the second being another and so on. Note that it is not defined behaviour if this list includes the joining peer, and it is up to you to ensure your network list does not contain duplicate peers. In all other cases the response will be a feedback message explaining any errors or results of other queries. You may notice that this is suspiciously similar to that of A3, therefore you can probably use some of your code from that if it helps.

The following error codes may be provided in response by the peer:

- 1 - OK (i.e. no problems encountered)
- 2 - Peer already exists (i.e. could not register a peer as they are already registered with the network)
- 3 - Bad Request (i.e. the request is coherent but cannot be

- servered as the file doesn't exist or is busy)
- 4 - Other (i.e. any error not covered by the other status codes)
 - 5 - Malformed (i.e. the request is malformed and could not be processed)

Submission

The submission should contain a file `src.zip` that contains the `src` directory of the handout; this should include any and all files necessary to run your code (including `csapp.c`, `csapp.h`, your `Makefile`, and any new files or test programs that you may have written).

Any Python code should also be included in your submission as it may form part of how you tested your C code (hint). As this course is not a Python course, you will not be marked according to the quality of your Python code.

Alongside the `src.zip` containing your code, submit a `report.pdf`, and a `group.txt`. `group.txt` must list the KU ids of your group members, one per line, and do so using *only* characters from the following set:

$$\{0x0A\} \cup \{0x30, 0x31, \dots, 0x39\} \cup \{0x61, 0x62, \dots, 0x7A\}$$

Please make sure your submission does not contain unnecessary files, including (but not limited to) compiled object files, binaries, or auxiliary files produced by editors or operating systems.