




Network layer: Routing

Michael Kirkedal Thomsen

Based on slides compiled by
Marcos Vaz Salles




Recap: Network layer

- What is the “best effort guarantee” of network layer ?
- Why can fragmentation happen at routers ?
How does IPv4 handle it ? Why does IPv6 not handle it ?
- What do forwarding tables in routers contain ?
Why is longest prefix match chosen ?
- Why is an IP address hierarchical ? What is a subnet mask ?
- How are IP addresses allocated ?



Forwarding Revisited

- How to resolve multiple matches?
 - Router identifies most specific prefix:
longest prefix match (LPM)
 - Cute algorithmic problem to achieve fast lookups

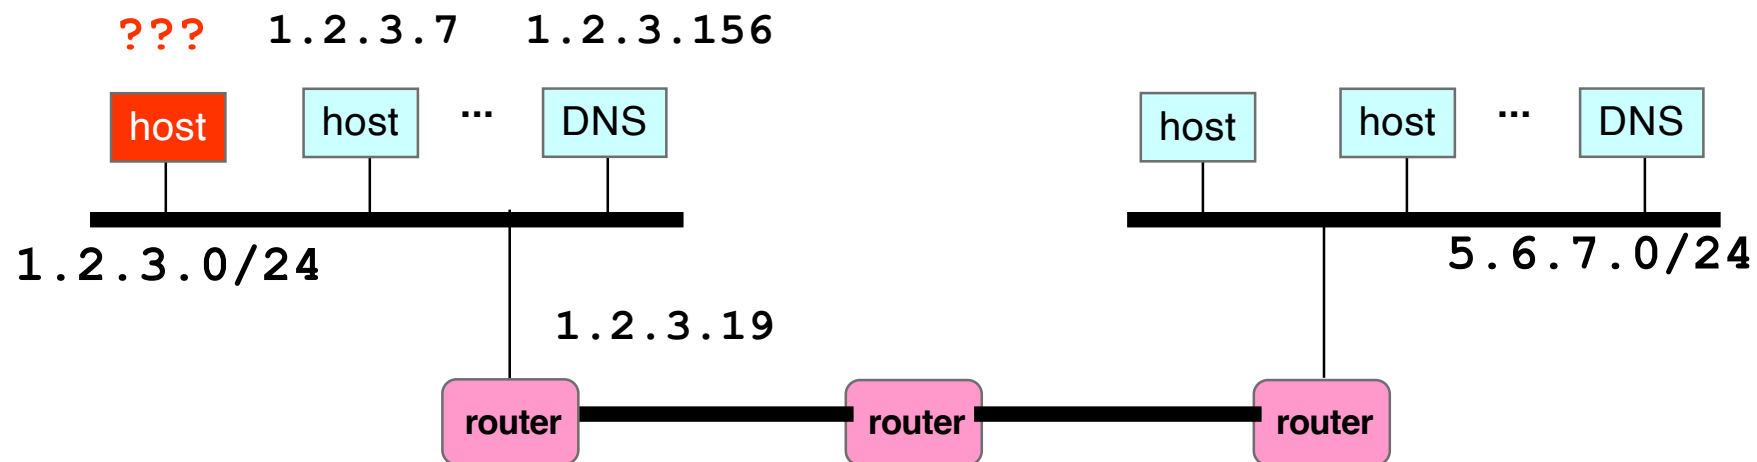


Source: Freedman (partial)



Recap: How To Bootstrap an End Host?


- What local Domain Name System server to use?
- What IP address the host should use?
- How to send packets to remote destinations?
- How to ensure incoming packets arrive?



Source: Freedman



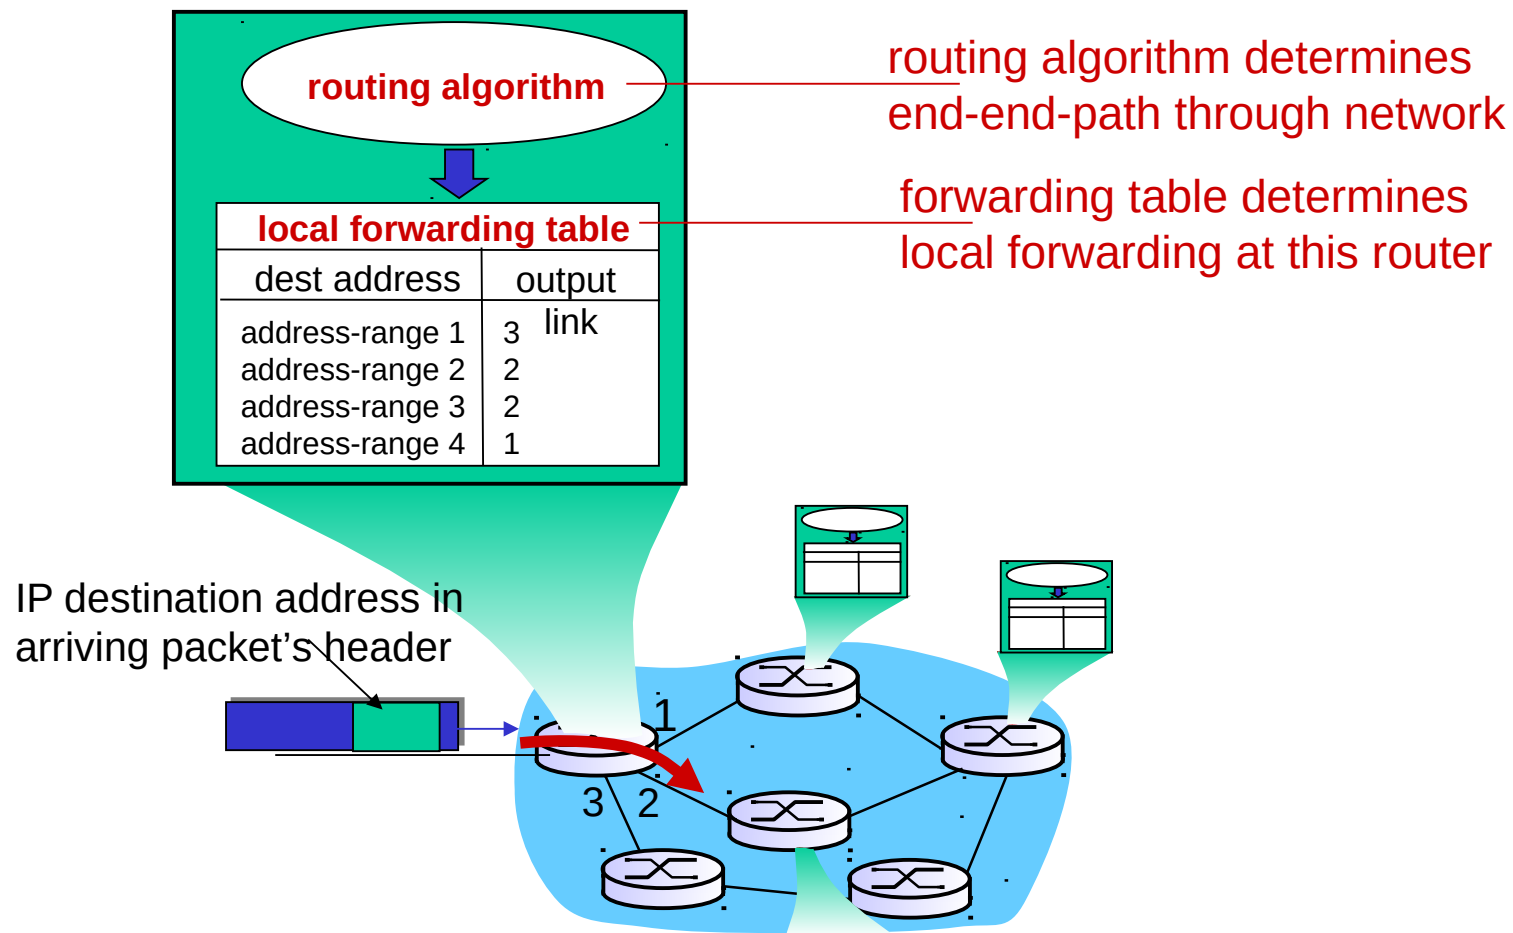
Routing: Mapping Link to Path



Source:
Freedman



Interplay of routing and forwarding



Three Issues to Address

- What does the protocol compute?
 - E.g., shortest paths
- What algorithm does the protocol run?
 - E.g., link-state routing
- How do routers learn end-host locations?
 - E.g., injecting into the routing protocol



Routing Algorithm Classification

Q: global or decentralized information?

global:

- all routers have complete topology, link cost info
- “link state” algorithms

decentralized:

- router knows physically-connected neighbors, link costs to neighbors
- iterative process of computation, exchange of info with neighbors
- “distance vector” algorithms

Q: static or dynamic?

static:

- routes change slowly over time

dynamic:


- routes change more quickly
 - periodic update
 - in response to link cost changes

Source: Kurose & Ross



What to Compute ?


- Shortest path(s) between pairs of nodes
 - A shortest-path tree rooted at each node
 - Min hop count or min sum of edge weights



Source: Freedman

Shortest Path Problem

- Compute: path costs to all nodes
 - From a given source u to all other nodes
 - Cost of the path through each outgoing link
 - Next hop along the least-cost path to s



Link State : Dijkstra's Algorithm

- Flood the topology information to all nodes
- Each node computes shortest paths to other nodes

Initialization

```
S = {u}
for all nodes v
  if (v is adjacent to
    u)
    D(v) = c(u,v)
  else D(v) = ∞
```


Loop

```
add w with smallest D(w) to S
update D(v) for all adjacent v:
  D(v) = min{D(v), D(w) +
    c(w,v)}
until all nodes are in S
```



Used in OSPF and IS-IS






Link State : Routing Example







Link State : Routing Example



Link State : Routing Example




Link State : Routing Example





Source: Freedman






Link State : Routing Example (contd.)



Link State : Routing Example (contd.)







Link State : Routing Example (contd.)



Source: Freedman



Link State : Routing Example (contd.)




Source: Freedman



Link State : Routing Example (contd.)

- Shortest-path tree from u
- Forwarding table at u



dest	link
v	(u,v)
w	(u,w)
x	(u,w)
y	(u,v)
z	(u,v)
s	(u,w)
t	(u,w)

Source: Freedman




Link State Algorithm Discussion

algorithm complexity: n nodes

- each iteration: need to check all nodes, w , not in N
- $n(n+1)/2$ comparisons: $O(n^2)$
- more efficient implementations possible: $O(n \log n)$


oscillations possible:

- e.g., support link cost equals amount of carried traffic:



Distance Vector : Bellman Ford Algorithm

- Define distances at each node x
 - $d_x(y)$ = cost of least-cost path from x to y
- Update distances based on neighbors
 - $d_x(y) = \min \{c(x,v) + d_v(y)\}$ over all neighbors v




$$d_u(z) = \min \{ c(u,v) + d_v(z), c(u,w) + d_w(z) \}$$

Used in RIP and EIGRP

Source: Freedman



Distance Vector : Routing Example




$$d_y(z) = 1$$

$$d_x(z) = 4$$




Distance Vector : Routing Example



$$d_y(z) = 1$$


$$d_x(z) = 4$$



$$d_v(z) = \min \{ 2 + d_y(z), 1 + d_x(z) \} = 3$$





Distance Vector : Routing Example (contd.)



$$\begin{aligned} d_w(z) &= \min\{1+d_x(z), \\ &\quad 4+d_s(z), \\ &\quad 2+d_u(z)\} \\ &= 5 \end{aligned}$$

Distance Vector : Routing Example (contd.)




$$d_w(z) = \min\{1+d_x(z), \\ 4+d_s(z), \\ 2+d_u(z)\} \\ = 5$$

$$d_u(z) = ??$$


(A) 5 (B) 6 (C) 7



Distance Vector : Routing Example (contd.)



$$d_w(z) = \min\{1+d_x(z), 4+d_s(z), 2+d_u(z)\} = 5$$




$$d_u(z) = \min\{3+d_v(z), 2+d_w(z)\} = 6$$



Distance Vector : Link Cost Changes


link cost changes:

- Node detects local link cost change
- Updates routing info, recalculates distance vector
- If DV changes, notify neighbors
- **Good News travels fast**



link cost changes:

- Node detects local link cost change
- **Bad news travels slow** - “count to infinity” problem!
- 44 iterations before algorithm stabilizes: see text
- **Poisoned Reverse for faster convergence.** Will this completely solve count to infinity problem?



Source: Kurose & Ross

Distance Vector : Link Cost Changes

message complexity

- **LS:** with n nodes, E links, $O(nE)$ msgs sent
- **DV:** exchange between neighbors only
 - convergence time varies

speed of convergence

- **LS:** $O(n^2)$ algorithm requires $O(nE)$ msgs
 - may have oscillations
- **DV:** convergence time varies
 - may be routing loops
 - count-to-infinity problem

robustness: what happens if router malfunctions?

LS:

- node can advertise incorrect *link* cost
- each node computes only its *own* table

DV:

- DV node can advertise incorrect *path* cost
- each node's table used by others
 - error propagate through network

Source: Kurose & Ross



Routing Issues

Our routing study thus far - idealization

- All routers identical
- Network “flat”
... *not* true in practice

Scale: with 600 million destinations:

- Can't store all dest's in routing tables!
- Routing table exchange would swamp links!

Administrative autonomy

- Internet = network of networks
- Each network admin may want to control routing in its own network

Source: Kurose & Ross



Hierarchical Routing – Standard CS trick

- Aggregate routers into regions, “autonomous systems” (AS)
- Routers in same AS run same routing protocol
 - “Intra-AS” routing protocol
 - Routers in different AS can run different intra-AS routing protocol

Gateway router:

- At “edge” of its own AS
- Has link to router in another AS

Source: Kurose & Ross



Summary

- Routing Algorithms are graph based
 - Centralized → Link State
 - Distributed → Distance Vector
- Routing algorithms have different characteristics
- Many other hacks too! 😊
 - Tunneling, firewalls, mobile gateways, VPNs

